

```
In [73]: #Load essential libraries for processing and visualising data
```

```
import numpy as np
import pandas as pd          #Processing data
import matplotlib.pyplot as plt #Visualisation
import seaborn as sns        #Visualisation
```

```
In [74]: df = pd.read_excel('NSW_Road_Crash_Data_2017-2021_CRASH.xlsx')
df.head()
```

```
Out[74]:
```

	CrashID	Degree_of_crash	Degree_of_crash_detailed	Reporting_year	Year_of_crash	Month_of_crash	Day_of_week_of_crash	Two_hour_intervals	Street_of_cra:
0	1122708	Fatal	Fatal	2017.0	2017.0	January	Monday	18:00 - 19:59	HOLLOWA'
1	1122709	Fatal	Fatal	2017.0	2017.0	January	Monday	12:00 - 13:59	PUT
2	1122710	Fatal	Fatal	2017.0	2017.0	January	Tuesday	14:00 - 15:59	IRRIGATIC
3	1123942	Fatal	Fatal	2017.0	2017.0	January	Thursday	10:00 - 11:59	VARD'
4	1123948	Fatal	Fatal	2017.0	2017.0	January	Saturday	12:00 - 13:59	PRINC

5 rows × 49 columns

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101718 entries, 0 to 101717
Data columns (total 49 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CrashID                               101718 non-null  int64
1   Degree_of_crash                       101718 non-null  object
2   Degree_of_crash_detailed              101714 non-null  object
3   Reporting_year                        101714 non-null  float64
4   Year_of_crash                         101714 non-null  float64
5   Month_of_crash                        101714 non-null  object
6   Day_of_week_of_crash                  101713 non-null  object
7   Two_hour_intervals                    101713 non-null  object
8   Street_of_crash                       101714 non-null  object
9   Street_type                           101714 non-null  object
10  Distance                              101715 non-null  float64
11  Direction                             101714 non-null  object
12  Identifying_feature                   101714 non-null  object
13  Identifying_feature_type              101718 non-null  object
14  Town                                  101717 non-null  object
15  Route_no                              64996 non-null  float64
16  School_zone_location                 101717 non-null  object
17  School_zone_active                   101717 non-null  object
18  Type_of_location                     101716 non-null  object
19  Latitude                             101715 non-null  float64
20  Longitude                             101715 non-null  float64
21  LGA                                   101715 non-null  object
22  Urbanisation                          101716 non-null  object
23  Alignment                             101716 non-null  object
24  Primary_permanent_feature             24733 non-null  object
25  Primary_temporary_feature             1487 non-null   object
26  Primary_hazardous_feature             2045 non-null   object
27  Street_lighting                       101717 non-null  object
28  Road_surface                          101717 non-null  object
29  Surface_condition                     101717 non-null  object
30  Weather                               101717 non-null  object
31  Natural_lighting                      101718 non-null  object
32  Signals_operation                     101718 non-null  object
33  Other_traffic_control                  101718 non-null  object
34  Speed_limit                           101718 non-null  object
35  Road_classification                   101718 non-null  object
36  RUM_code                              101718 non-null  int64
37  RUM_description                       101718 non-null  object
38  DCA_code                              101718 non-null  int64
39  DCA_description                       101718 non-null  object
40  DCA_supplement                        12695 non-null  object
41  First_impact_type                     101718 non-null  object
42  Key_TU_type                           101718 non-null  object
43  Other_TU_type                          72060 non-null  object
44  No_of_traffic_units_involved           101718 non-null  int64
45  No_killed                             101718 non-null  int64
46  No_seriously_injured                  101718 non-null  int64
47  No_moderately_injured                 101718 non-null  int64
48  No_minor_other_injured                101718 non-null  int64
dtypes: float64(6), int64(8), object(35)
memory usage: 38.0+ MB
```

Data Cleaning

```
In [84]: #Categorical columns
cat_cols = df.select_dtypes(include=['object', 'category']).columns
```

```
cat_cols
```

```
Out[84]: Index(['Degree_of_crash', 'Degree_of_crash_detailed', 'Month_of_crash',
        'Day_of_week_of_crash', 'Two_hour_intervals', 'Street_of_crash',
        'Street_type', 'Direction', 'Identifying_feature',
        'Identifying_feature_type', 'Town', 'School_zone_location',
        'School_zone_active', 'Type_of_location', 'LGA', 'Urbanisation',
        'Alignment', 'Primary_permanent_feature', 'Primary_temporary_feature',
        'Primary_hazardous_feature', 'Street_lighting', 'Road_surface',
        'Surface_condition', 'Weather', 'Natural_lighting', 'Signals_operation',
        'Other_traffic_control', 'Speed_limit', 'Road_classification',
        'RUM_description', 'DCA_description', 'DCA_supplement',
        'First_impact_type', 'Key_TU_type', 'Other_TU_type'],
        dtype='object')
```

```
In [85]: #Replace null values in categorical columns with the mode
df[cat_cols] = df[cat_cols].apply(lambda x: x.fillna(x.mode()[0]))

print(df.isnull().sum())
```

```
CrashID                0
Degree_of_crash        0
Degree_of_crash_detailed 0
Reporting_year         4
Year_of_crash          4
Month_of_crash         0
Day_of_week_of_crash   0
Two_hour_intervals     0
Street_of_crash        0
Street_type            0
Distance              3
Direction              0
Identifying_feature     0
Identifying_feature_type 0
Town                  0
Route_no              36722
School_zone_location   0
School_zone_active     0
Type_of_location       0
Latitude              3
Longitude             3
LGA                   0
Urbanisation           0
Alignment              0
Primary_permanent_feature 0
Primary_temporary_feature 0
Primary_hazardous_feature 0
Street_lighting        0
Road_surface           0
Surface_condition      0
Weather                0
Natural_lighting       0
Signals_operation      0
Other_traffic_control   0
Speed_limit            0
Road_classification    0
RUM_code               0
RUM_description        0
DCA_code               0
DCA_description        0
DCA_supplement         0
First_impact_type      0
Key_TU_type            0
Other_TU_type          0
No_of_traffic_units_involved 0
No_killed              0
No_seriously_injured   0
No_moderately_injured  0
No_minor_other_injured 0
Crash_severe           0
dtype: int64
```

```
In [86]: #Numerical_columns
numerical_cols = df.select_dtypes(include=['float64', 'int']).columns
numerical_cols
```

```
Out[86]: Index(['CrashID', 'Reporting_year', 'Year_of_crash', 'Distance', 'Route_no',
        'Latitude', 'Longitude', 'RUM_code', 'DCA_code',
        'No_of_traffic_units_involved', 'No_killed', 'No_seriously_injured',
        'No_moderately_injured', 'No_minor_other_injured', 'Crash_severe'],
        dtype='object')
```

```
In [87]: #Replace null values in numerical columns with the mean
df[numerical_cols] = df[numerical_cols].apply(lambda x: x.fillna(x.mean()))

print(df.isnull().sum())
```

```

CrashID                0
Degree_of_crash        0
Degree_of_crash_detailed  0
Reporting_year         0
Year_of_crash          0
Month_of_crash         0
Day_of_week_of_crash   0
Two_hour_intervals     0
Street_of_crash        0
Street_type            0
Distance              0
Direction             0
Identifying_feature    0
Identifying_feature_type 0
Town                  0
Route_no              0
School_zone_location   0
School_zone_active     0
Type_of_location       0
Latitude              0
Longitude             0
LGA                   0
Urbanisation           0
Alignment             0
Primary_permanent_feature 0
Primary_temporary_feature 0
Primary_hazardous_feature 0
Street_lighting        0
Road_surface           0
Surface_condition      0
Weather               0
Natural_lighting       0
Signals_operation      0
Other_traffic_control  0
Speed_limit            0
Road_classification    0
RUM_code              0
RUM_description        0
DCA_code              0
DCA_description        0
DCA_supplement         0
First_impact_type      0
Key_TU_type           0
Other_TU_type          0
No_of_traffic_units_involved 0
No_killed              0
No_seriously_injured   0
No_moderately_injured  0
No_minor_other_injured 0
Crash_severe           0
dtype: int64

```

```
In [88]: df.duplicated().sum() #Check duplicates
```

```
Out[88]: 0
```

Visualisation

Bar chart - top five towns/suburbs with highest no. of crashes

```

In [15]: top_10_towns = df.groupby('Town')['CrashID'].count().sort_values(ascending=False).head(10) #Select 10 towns with highest accident/cra
top_10_towns_reorder = top_10_towns.sort_values(ascending=True) #order them in ascending order

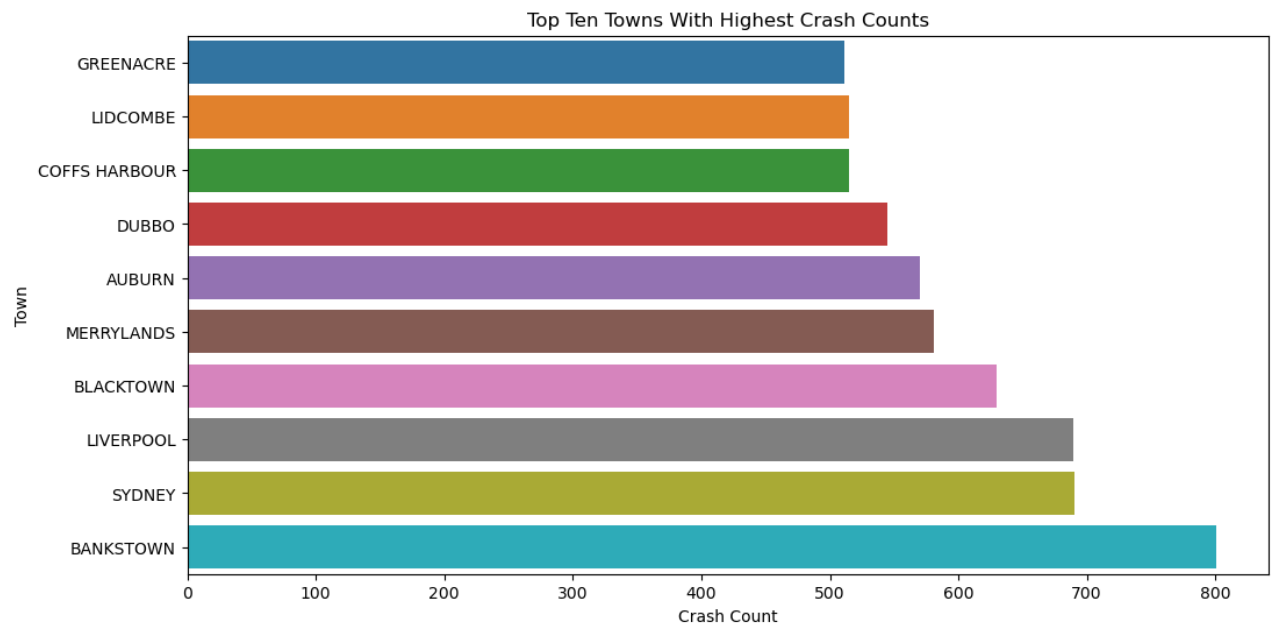
plt.figure(figsize=(12, 6)) #fix the visual size

sns.barplot(x=top_10_towns_reorder.values, y=top_10_towns_reorder.index, orient='h') #seaborn horizontal bar plot

plt.xlabel('Crash Count')
plt.ylabel('Town')
plt.title('Top Ten Towns With Highest Crash Counts')

plt.show()

```



Pie chart - display the proportion of fatalities occurring in different weathers

```
In [22]: plt.style.use('classic')
plt.style.use('default')

# Filter out data from the year 2016
df1 = df[df['Reporting_year'] != 2016]

# Filter out Unknown weather
df1 = df1[df1['Weather'] != 'Unknown']
df1 = df1[df1['Weather'] != 'Other']

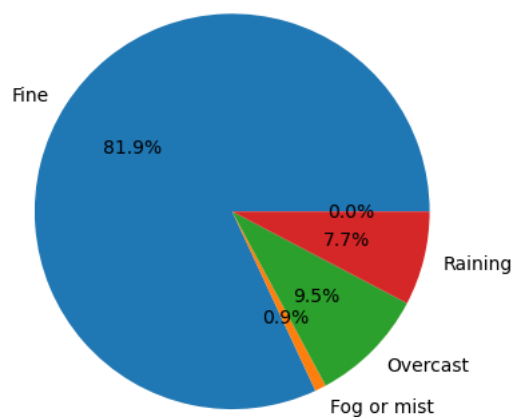
# Remove duplicates based on specific columns
d1 = df1.drop_duplicates(subset=['CrashID'])

plt.figure(figsize=(12, 6))
df1.groupby('Weather').sum().plot(kind='pie',
                                  title='Distribution of fatalities by weather conditions',
                                  y='No_killed', legend=False,
                                  ylabel='', autopct='%1.1f%%')

plt.figure(figsize=(12, 6)) #fix the visual size
plt.show()
```

<Figure size 1200x600 with 0 Axes>

Distribution of fatalities by weather conditions



<Figure size 1200x600 with 0 Axes>

Line plot & bar chart - distribution of the number of people killed on the road vary by month & year

```
In [23]: #create new dataframe for grouped data
grouped_data = df.groupby(['Year_of_crash', 'Month_of_crash'])['No_killed'].sum()
grouped_data = grouped_data.reset_index()

#Clean out rows where the year is 2016 or 0
grouped_data['Year_of_crash'] = grouped_data['Year_of_crash'].fillna(0).astype(int)
grouped_data = grouped_data[(grouped_data['Year_of_crash'] != 2016) & (grouped_data['Year_of_crash'] != 0)]
```

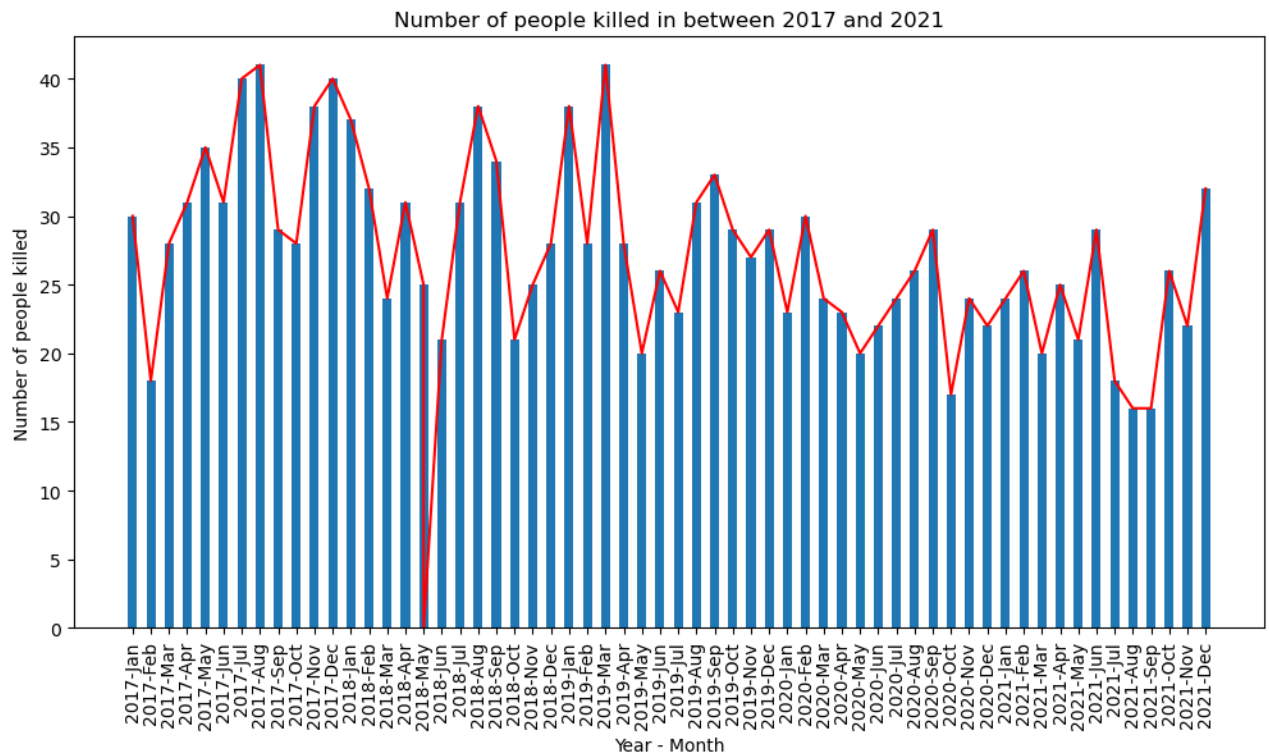
```
#creating format for x axis
grouped_data['Year-Month'] = grouped_data['Year_of_crash'].astype(str).str[:4] + '-' + grouped_data['Month_of_crash'].str[:3]

#map out the months using a dictionary to order it Later
month_dict = {'January': 1, 'February': 2, 'March': 3, 'April': 4, 'May': 5, 'June': 6, 'July': 7, 'August': 8, 'September': 9, 'October': 10, 'November': 11, 'December': 12}
grouped_data['Month_Number'] = grouped_data['Month_of_crash'].map(month_dict)

#Sort the data by the 'Year and Month_Number' column
grouped_data = grouped_data.sort_values(['Year_of_crash', 'Month_Number'])

#creating line
plt.figure(figsize=(12, 6))
plt.plot(grouped_data['Year-Month'], grouped_data['No_killed'], color = 'r')

#creating bar
plt.bar(grouped_data['Year-Month'], grouped_data['No_killed'], width=0.5)
plt.xticks(rotation = 90)
plt.xlabel("Year - Month")
plt.ylabel("Number of people killed")
plt.title("Number of people killed in between 2017 and 2021")
plt.show()
```



Line plot & bar chart - show the trend in the number of fatalities and injuries over years

```
In [25]: df['Year_of_crash'] = df['Year_of_crash'].astype(int)
#change dtype of year_of_crash column
df.dtypes #check if the dtype changed

#Total fatalities over year
year_total_fatal = df.groupby('Year_of_crash')['No_killed'].sum().reset_index()
year_total_fatal = year_total_fatal[(year_total_fatal['Year_of_crash'] != 0) & (year_total_fatal['Year_of_crash'] != 2016)] #remove 2016
year_total_fatal

#Total injuries over year
year_total_injury = df.groupby('Year_of_crash')[['No_seriously_injured', 'No_moderately_injured', 'No_minor_other_injured']].sum().reset_index()
year_total_injury['Total_Injuries'] = year_total_injury[['No_seriously_injured', 'No_moderately_injured', 'No_minor_other_injured']].sum(axis=1)
year_total_injury = year_total_injury[(year_total_injury['Year_of_crash'] != 0) & (year_total_injury['Year_of_crash'] != 2016)] #remove 2016
year_total_injury

import matplotlib.pyplot as plt #import visualisation library

year = year_total_fatal['Year_of_crash'] #select values from year_of_crash column in year_total_fatal dataframe for plotting
total_injuries = year_total_injury['Total_Injuries'] #select values from total_injuries column in year_total_fatal dataframe for plotting
total_fatalities = year_total_fatal['No_killed'] #select values from no_killed column in year_total_fatal dataframe for plotting

plt.figure(figsize=(12,6)) #define plot size
fig, ax1 = plt.subplots()

ax1.set_xlabel('Year')

#Create line plot for total fatalities --- total fatalities on the left y-axis
ax1.set_ylabel('Fatalities', color='black')
ax1.plot(year, total_fatalities, marker='o', color = 'red', label='Fatalities')

#Create barplot for total injuries -- total injuries on the right y-axis
```

```

ax2 = ax1.twinx()
ax2.set_ylabel('Injuries', color='black')
ax2.bar(year, total_injuries, color = 'blue', alpha = 0.2, label='Injuries')

# Display the fatality numbers
for i,j in zip(year, total_fatalities):
    ax1.annotate(str(j),xy=(i,j))

plt.title("Trends in Fatalities and Injuries Over the Years")


# Add a single legend for both bar and line plot
lines, labels = ax1.get_legend_handles_labels()
lines2, labels2 = ax2.get_legend_handles_labels()
lines = lines + lines2
labels = labels + labels2

# Display the combined legend in the upper right
ax2.legend(lines, labels, loc = 'upper right')

#Show the plot
plt.show()

```

<Figure size 1200x600 with 0 Axes>



Stacked bar chart - relationship between the speed limit and accident statistics

```

In [26]: # Define the desired order of speed limits
speed_limit_order = [
    '10 km/h', '20 km/h', '30 km/h',
    '40 km/h', '50 km/h', '60 km/h',
    '70 km/h', '80 km/h', '90 km/h',
    '100 km/h', '110 km/h'
]

# Convert 'Speed_Limit' to categorical with the desired order
df['Speed_limit'] = pd.Categorical(
    df['Speed_limit'], categories=speed_limit_order, ordered=True)

# Group the data by 'Speed_Limit' and 'Degree_of_crash' and count occurrences
grouped_data = df.groupby(['Speed_limit', 'Degree_of_crash']).size().unstack()

# Create a stacked bar chart
ax = grouped_data.plot(kind='bar', stacked=True, figsize=(12, 6), alpha=0.7) # Adjust alpha for visibility

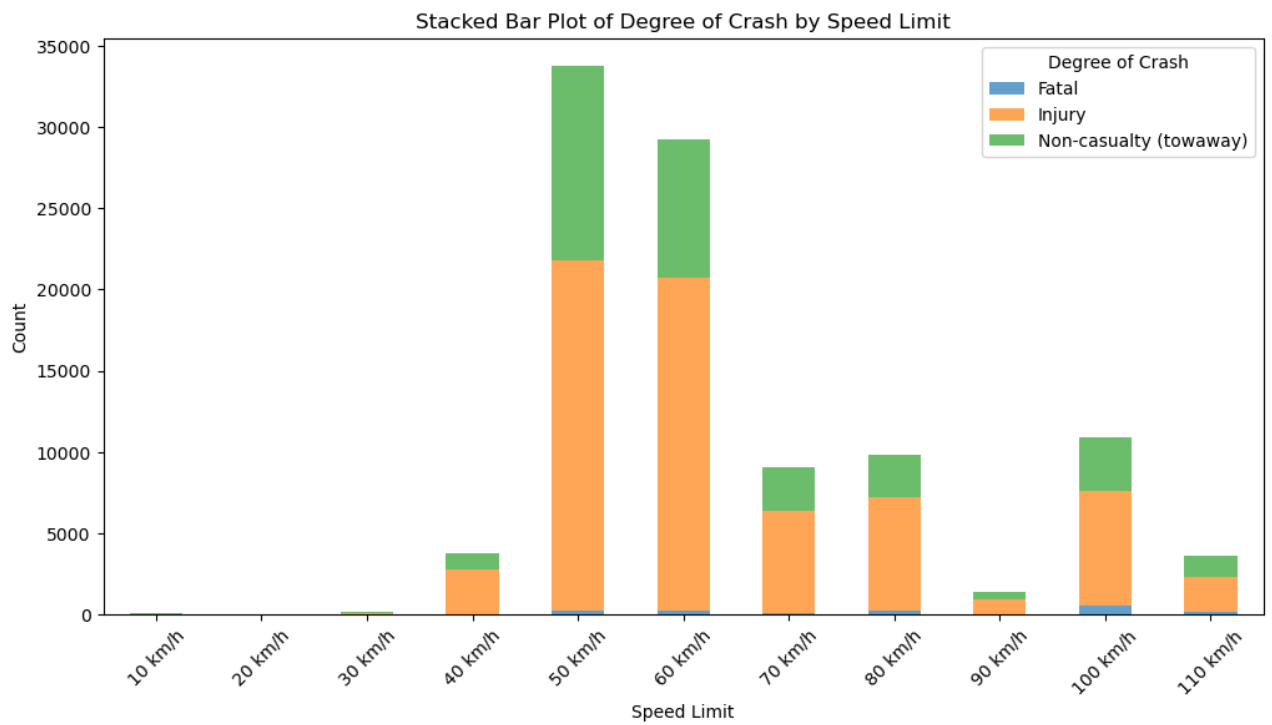
# Set labels and title
plt.xlabel("Speed Limit")
plt.ylabel("Count")
plt.title("Stacked Bar Plot of Degree of Crash by Speed Limit")

# Show the Legend
plt.legend(title="Degree of Crash")

# Rotate x-axis labels for better readability
ax.set_xticklabels(ax.get_xticklabels(), rotation=45)

# Show the plot
plt.show()

```



Bar chart - show most notorious locations in the past 5 years

```
In [28]: df6 = pd.DataFrame()
df6['Town'] = df['Town']
df6['Street_of_crash'] = df['Street_of_crash']
df6['Street_type'] = df['Street_type']
df6['Year_of_crash'] = df['Year_of_crash']

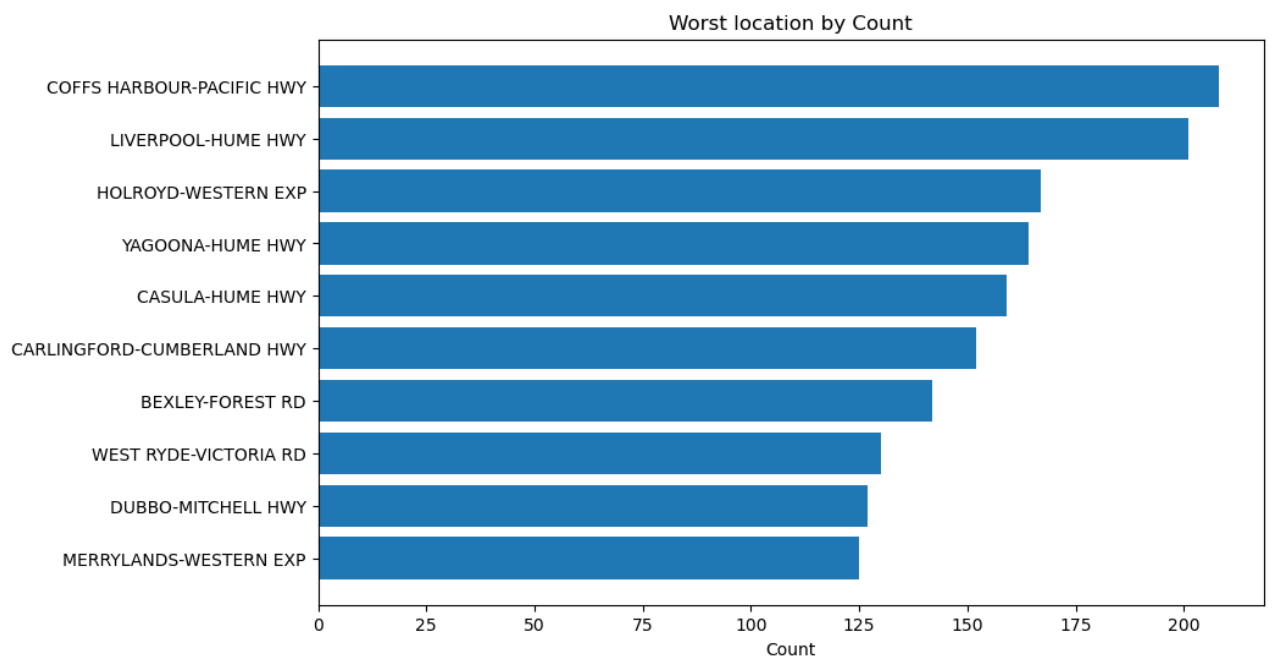
#Concat columns
df6['Concatenated'] = df6['Town'] + '-' + df6['Street_of_crash'] + ' ' + df6['Street_type']

#Clean out rows where the year is 2016 or 0
df6['Year_of_crash'] = df6['Year_of_crash'].astype(int)
df6 = df6[(df6['Year_of_crash'] != 2016) & (df6['Year_of_crash'] != 0)]

#Group by 'Concatenated' and count the number of occurrences
grouped_df6 = df6.groupby('Concatenated').size().reset_index(name='count')

#Select the top 10 locations
grouped_df6 = grouped_df6.sort_values('count', ascending=True).tail(10)

#Create a horizontal bar chart
plt.figure(figsize=(10, 6))
plt.barh(grouped_df6['Concatenated'], grouped_df6['count'])
plt.xlabel('Count')
plt.title('Worst location by Count')
plt.show()
```



Density heatmap - show crashes in geographical locations

```
In [30]: pip install dash #Install dash Library
```

Note: you may need to restart the kernel to use updated packages.

ERROR: Invalid requirement: '#Install'

```
In [31]: import plotly.express as px #import plotly package
```

```
In [33]: df7 = pd.DataFrame()
df7['Magnitude'] = df['No_killed'] + df['No_seriously_injured']
df7['Longitude'] = df['Longitude']
df7['Latitude'] = df['Latitude']
df7['Year_of_crash'] = df['Year_of_crash']

#Change data type for column year of crash
df7['Year_of_crash'] = df7['Year_of_crash'].astype(int)

#Clean out rows where the year is 2016 or 0
df7 = df7[(df7['Year_of_crash'] != 2016) & (df7['Year_of_crash'] != 0)]


#removes Magnitude where it is 0
df7 = df7[df7['Magnitude'] != 0]

#sets Limits
df7['Magnitude'] = df7['Magnitude'].clip(1, 8)

#creating density heatmap
fig = px.density_mapbox(df7, lat='Latitude', lon='Longitude', z='Magnitude', radius=10, center=dict(lat=0, lon=180), zoom=0, mapbox_s
fig.show()
```

Heatmap to observe the correlation between variables

```
In [34]: df_cleaned = df.drop(columns=['DCA_code', 'RUM_code'])
fig, ax = plt.subplots(figsize=(10,10))
cm = sns.heatmap(df_cleaned.corr(), linewidths = .5, cmap="YlGnBu", annot=True, ax=ax, fmt='.1g')
```

Predictive Modelling

Linear Regression - predict severity of crashes from speed limit

In [43]: `df.dtypes`

```
Out[43]: CrashID                int64
Degree_of_crash              object
Degree_of_crash_detailed      object
Reporting_year                int32
Year_of_crash                 int32
Month_of_crash                object
Day_of_week_of_crash          object
Two_hour_intervals            object
Street_of_crash               object
Street_type                   object
Distance                      float64
Direction                     object
Identifying_feature           object
Identifying_feature_type      object
Town                          object
Route_no                      float64
School_zone_location          object
School_zone_active            object
Type_of_location              object
Latitude                      float64
Longitude                     float64
LGA                           object
Urbanisation                  object
Alignment                     object
Primary_permanent_feature      object
Primary_temporary_feature      object
Primary_hazardous_feature      object
Street_lighting                object
Road_surface                   object
Surface_condition              object
Weather                        object
Natural_lighting               object
Signals_operation              object
Other_traffic_control          object
Speed_limit                    category
Road_classification            object
RUM_code                       int64
RUM_description                object
DCA_code                       int64
DCA_description                object
DCA_supplement                 object
First_impact_type              object
Key_TU_type                    object
Other_TU_type                  object
No_of_traffic_units_involved   int64
No_killed                      int64
No_seriously_injured           int64
No_moderately_injured          int64
No_minor_other_injured         int64
dtype: object
```

```
In [50]: #Preparing data for modelling

#df['Year_of_crash'] = df['Year_of_crash'].fillna(0).astype(int)
#df['Reporting_year'] = df['Reporting_year'].fillna(0).astype(int)
df = df[(df['Year_of_crash'] != 2016) & (df['Year_of_crash'] != 0)]
df = df[(df['Reporting_year'] != 2016) & (df['Reporting_year'] != 0)]

df_data = df[['Speed_limit', 'No_of_traffic_units_involved', 'No_killed', 'No_seriously_injured', 'No_moderately_injured', 'No_minor_o

#df_data.columns[df_data.isna().any()].tolist()

df_data = df_data[df_data['Speed_limit'] != 'Unknown'].copy()
df_data['Speed_limit'].unique()

#df_data.head()

df_data['Speed_limit'] = df_data['Speed_limit'].str.extract('(\d+)').fillna(0).astype(int)

# Add the Crash_Severity column to the DataFrame
df_data['Crash_Severity'] = (
    df_data['No_killed'] * 16 +
    df_data['No_seriously_injured'] * 8 +
    df_data['No_moderately_injured'] * 4 +
    df_data['No_minor_other_injured'] * 2 +
    df_data['No_of_traffic_units_involved'] * 1
)
```

```
In [51]: #Modelling Regression model

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

# Extract features and target variable
X = df_data[['Speed_limit']]
y = df_data['Crash_Severity']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a Linear regression model
model = LinearRegression()

# Train the model
model.fit(X_train, y_train)
```

```

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

plt.figure(figsize=(10, 6))
sns.boxplot(x='Speed_limit', y='Crash_Severity', data=df_data)
plt.title('Crash Severity Distribution for Different Speed Limits')
plt.show()

print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')

print(f'Coefficient: {model.coef_[0]}')
print(f'Intercept: {model.intercept_}')

import matplotlib.pyplot as plt


# Plot the data points
plt.scatter(X_test, y_test, color='black', label='Data', s=5)

# Plot the linear regression line
plt.plot(X_test, y_pred, color='blue', linewidth=1, label='Linear Regression Line')

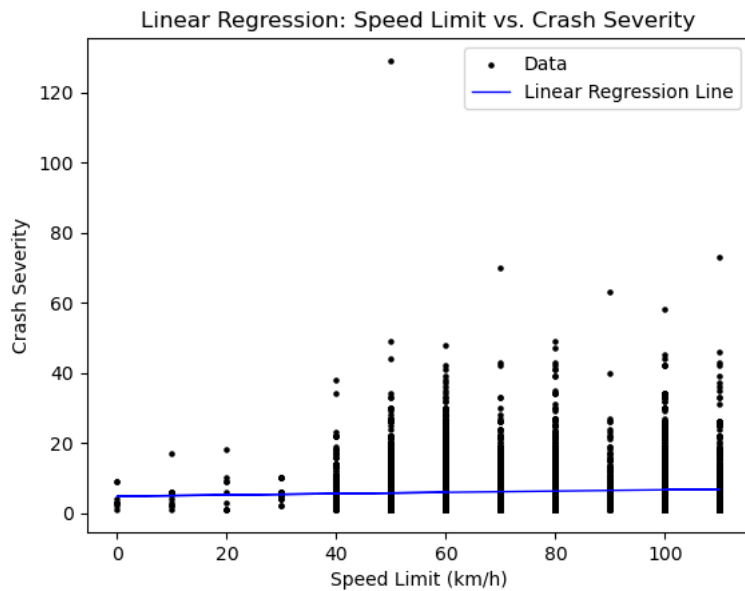
# Labeling the plot
plt.title('Linear Regression: Speed Limit vs. Crash Severity')
plt.xlabel('Speed Limit (km/h)')
plt.ylabel('Crash Severity')
plt.legend()

# Show the plot
plt.show()

```



Mean Squared Error: 23.699084817856665
 R-squared: 0.006749663582501952
 Coefficient: 0.018437199801925434
 Intercept: 4.755657655367552



Decision Tree - improve model; predict severity of crashes based on speed limit

```
In [53]: #Preparing data for modelling decision tree

#convert year to integers
#df['Year_of_crash'] = df['Year_of_crash'].fillna(0).astype(int)
#df['Reporting_year'] = df['Reporting_year'].fillna(0).astype(int)

#exclude 2016
#df = df[(df['Year_of_crash'] != 2016) & (df['Year_of_crash'] != 0)]
#df = df[(df['Reporting_year'] != 2016) & (df['Reporting_year'] != 0)]

#Select columns
df_data = df[['Degree_of_crash_detailed', 'Speed_limit']].copy()

df_data.columns[df_data.isna().any()].tolist()

#Remove unknown from speed limit
df_data = df_data[df_data['Speed_limit'] != 'Unknown'].copy()
df_data['Speed_limit'].unique()

df_data.head()

#Convert string into integer, extracting number
df_data['Speed_limit'] = df_data['Speed_limit'].str.extract('(\d+)').fillna(0).astype(int)
df_data.head()
```

```
Out[53]:
```

	Degree_of_crash_detailed	Speed_limit
0	Fatal	50
1	Fatal	80
2	Fatal	100
3	Fatal	60
4	Fatal	100

```
In [54]: #Decision tree modelling

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.tree import plot_tree
from sklearn.metrics import confusion_matrix

#Split data
X = df_data[['Speed_limit']]
y = df_data['Degree_of_crash_detailed']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = DecisionTreeClassifier(random_state=42)

#Train model
model.fit(X_train, y_train)

#Make prediction
y_pred = model.predict(X_test)

#Print report
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')
```

```
print('\nClassification Report:')
print(classification_report(y_test, y_pred))

#Visualise tree
plt.figure(figsize=(15, 8))
plot_tree(model, filled=True, feature_names=['Speed_limit'], class_names=None)
plt.show()
```

Accuracy: 0.3206294566019179

Classification Report:

C:\Users\maikh\anaconda3\lib\site-packages\sklearn\metrics_classification.py:1318: UndefinedMetricWarning:

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

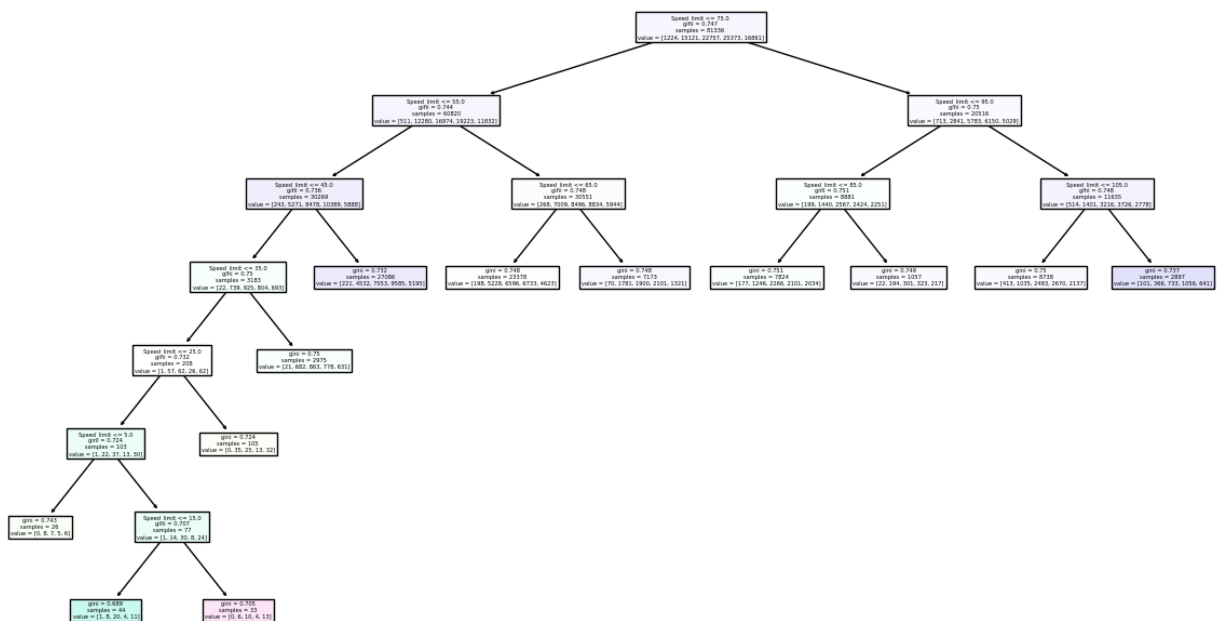
C:\Users\maikh\anaconda3\lib\site-packages\sklearn\metrics_classification.py:1318: UndefinedMetricWarning:

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

C:\Users\maikh\anaconda3\lib\site-packages\sklearn\metrics_classification.py:1318: UndefinedMetricWarning:

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

	precision	recall	f1-score	support
Fatal	0.00	0.00	0.00	306
Minor/Other Injury	0.22	0.00	0.00	3838
Moderate Injury	0.30	0.15	0.20	5629
Non-casualty (towaway)	0.32	0.89	0.47	6353
Serious Injury	0.40	0.00	0.00	4209
accuracy			0.32	20335
macro avg	0.25	0.21	0.14	20335
weighted avg	0.31	0.32	0.20	20335



Linear Regression - Severity of crash vs. Day of week

In [57]: `df_no_nan = df.dropna(subset=['Degree_of_crash', 'No_killed', 'Day_of_week_of_crash'])`

```
from sklearn import linear_model

#clarifying each category into numerical values
def crash_severity(value):
    if value == "Fatal":
        return 5
    elif value == "Injury":
        return 3
    else:
        return 1
#assigning numerical values to days of week
def weekdayno(value):
    if value == "Monday":
        return 1
    elif value == "Tuesday":
        return 2
    elif value == "Wednesday":
        return 3
    elif value == "Thursday":
        return 4
```

```

elif value == "Friday":
    return 5
elif value == "Saturday":
    return 6
elif value == "Sunday":
    return 7

df_no_nan['weekdayno'] = df_no_nan['Day_of_week_of_crash'].map(weekdayno)
df_no_nan['crash_severity'] = df_no_nan['Degree_of_crash'].map(crash_severity)

#df_no_nan.head()

regr = linear_model.LinearRegression()

severity= df_no_nan[['crash_severity']]
weekdayno= df_no_nan[['weekdayno']]

regr.fit(weekdayno, severity)

plt.scatter( df_no_nan.weekdayno, df_no_nan.crash_severity, color='Black')

plt.plot(weekdayno, regr.coef_[0][0]*weekdayno + regr.intercept_[0], '-r')
plt.xlabel("weekday")
plt.ylabel("crash severity")
plt.title('crash severity based on day of week', fontsize=14)

```

Out[57]: Text(0.5, 1.0, 'crash severity based on day of week')



Logistic Regression - predict the likelihood of a crash being severe based on features like weather, road alignment, and road classification

```

In [75]: #Import important Libraries
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression

```

```

In [77]: #Step 1 - Separate x and y data
from sklearn.preprocessing import LabelEncoder

X1 = df[['Weather', 'Alignment', 'Road_classification']] #predictor variables
print(X1)

#Because the Logistic reg only accepts categorical or numerical independent variables
X1 = pd.get_dummies(X1, columns=['Weather', 'Alignment', 'Road_classification'])

#Understanding severity of crashes
#road_accident['Degree_of_crash'].value_counts()

#There are 3 types of crash outcomes
#In order to conduct a Logistic regression, these outcomes must be
#categorized into binary variables
#Decision: map fatal to value 1, Injury & Towaway to value 0

#Create binary variables to the crash severity variable
df['Crash_severe'] = df['Degree_of_crash'].map({'Fatal':1, 'Injury':1, 'Non-casualty (towaway)':0})
y1 = df[['Crash_severe']]
print(y1)

#Step 2 - Training and testing set split

```

```
X1_train, X1_test, y1_train, y1_test = train_test_split(
    X1, y1, test_size=0.2, random_state=42, stratify=y1
)
```

#After this, use resampling method: undersampling

```

      Weather Alignment Road_classification
0      Overcast      Curved              Local
1      Overcast      Curved              State
2          Fine      Straight              State
3      Overcast      Straight          Regional
4          Fine      Curved              State
...      ...      ...      ...
101713      Fine      Curved              State
101714      Fine      Curved          Regional
101715      Overcast      Straight          Local
101716      Fine      Straight              State
101717      Fine      Straight          Local

```

[101718 rows x 3 columns]

Crash_severe

```

0          1
1          1
2          1
3          1
4          1
...      ...
101713      1
101714      1
101715      0
101716      0
101717      1

```

[101718 rows x 1 columns]

In [78]: `pip install -U imbalanced-learn #Import package for undersampling`

Note: you may need to restart the kernel to use updated packages.

ERROR: Invalid requirement: '#Import'

In [79]: *#Under-resampling technique*

```

from imblearn.under_sampling import RandomUnderSampler # Import the necessary Libraries

rus = RandomUnderSampler(random_state=42, sampling_strategy = 'majority')

# Balancing the data
X_resampled, y_resampled = rus.fit_resample(X1_train, y1_train)

```

In [81]: *#Step 3 - build Logistic regression model*
`r_accident_lr = LogisticRegression(solver = 'liblinear', random_state = 42, C = 10)`

#Step 4 - Fit the model to the resampled data
`r_accident_lr.fit(X_resampled, y_resampled)`

#Step 5 - Predict the test data

`y1_pred = r_accident_lr.predict(X1_test)`

`print(pd.DataFrame(y1_pred))`

#Step 6 - Evaluate test the model
`r_accident_lr.score(X1_test, y1_test)`

#Evaluation using confusion matrix
`from sklearn.metrics import confusion_matrix`
`import matplotlib.pyplot as plt`
`import seaborn as sns`

`def confusion_matrix_visual(y1_true, y1_pred, labels):`
 `cm = confusion_matrix(y1_true, y1_pred)`

```

    plt.figure(figsize=(6, 3))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.title('Confusion Matrix')
    plt.show()

```

`confusion_matrix_visual(y1_test, y1_pred, ['low', 'high'])`

#Precision report
`from sklearn.metrics import classification_report`

`print("Report : ",`
 `classification_report(y1_test, y1_pred))`

ROC Curves
`from sklearn.metrics import roc_curve, auc`
`import matplotlib.pyplot as plt`

`fpr, tpr, thresholds = roc_curve(y1_test, r_accident_lr.predict_proba(X1_test)[:, 1])`

```

roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='red', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()

```

C:\Users\maikh\anaconda3\lib\site-packages\sklearn\utils\validation.py:993: DataConversionWarning:

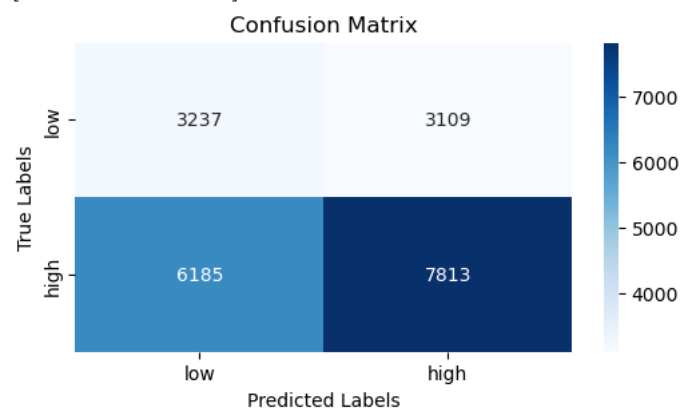
A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```


0
0 1
1 1
2 1
3 0
4 0
... ..
20339 0
20340 1
20341 1
20342 1
20343 1

```

[20344 rows x 1 columns]



Report :		precision	recall	f1-score	support
	0	0.34	0.51	0.41	6346
	1	0.72	0.56	0.63	13998
	accuracy			0.54	20344
	macro avg	0.53	0.53	0.52	20344
	weighted avg	0.60	0.54	0.56	20344





```
In [82]: # Precision-Recall Curves
from sklearn.metrics import precision_recall_curve, auc
import matplotlib.pyplot as plt

# Calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(y1_test, r_accident_lr.predict_proba(X1_test)[: , 1])

# Calculate AUC for precision-recall curve
pr_auc = auc(recall, precision)

# Plot precision-recall curve
plt.figure(figsize=(8, 6))
plt.plot(recall, precision, color='red', lw=2, label='PR curve (area = %0.2f)' % pr_auc)
plt.axhline(y=y1_test.mean().iloc[0], color='blue', linestyle='--', label='Baseline')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.legend(loc="lower left")
plt.show()
```



Feature selection for predictive models - using Chi-square method

```
In [89]: import pandas as pd
from sklearn.feature_selection import chi2
from sklearn.preprocessing import LabelEncoder

# Convert categorical variables to numerical Labels using LabelEncoder
label_encoder = LabelEncoder()
encoded_data = df.copy()
categorical_columns = ['Month_of_crash', 'Day_of_week_of_crash', 'Two_hour_intervals', 'Street_of_crash', 'Street_type', 'Direction',

for col in categorical_columns:
    encoded_data[col] = label_encoder.fit_transform(df[col])

# Create a DataFrame with only the relevant columns
X_categorical = encoded_data[categorical_columns]
y = encoded_data['Degree_of_crash']

# Apply chi-squared test
chi2_stat, p_values = chi2(X_categorical, y)

# Create a DataFrame to display the results
chi2_results = pd.DataFrame({
    'Feature': categorical_columns,
    'Chi2 Statistic': chi2_stat,
    'P-Value': p_values
})

# Display the results
print(chi2_results)
```

	Feature	Chi2 Statistic	P-Value
0	Month_of_crash	9.010312	1.105187e-02
1	Day_of_week_of_crash	17.417811	1.651089e-04
2	Two_hour_intervals	27.484933	1.075778e-06
3	Street_of_crash	24857.818780	0.000000e+00
4	Street_type	255.626863	3.099788e-56
5	Direction	0.533835	7.657364e-01
6	Identifying_feature	102202.978255	0.000000e+00
7	Identifying_feature_type	835.101947	4.570000e-182
8	Town	199.402831	5.014475e-44
9	School_zone_location	40.797710	1.383215e-09
10	School_zone_active	0.990990	6.092691e-01
11	Type_of_location	2063.758680	0.000000e+00
12	LGA	969.804882	2.567708e-211
13	Urbanisation	695.568245	9.104889e-152
14	Alignment	78.515515	8.924265e-18
15	Primary_permanent_feature	675.994832	1.620270e-147
16	Primary_temporary_feature	0.429538	8.067278e-01
17	Primary_hazardous_feature	26.674858	1.612977e-06
18	Street_lighting	12371.234424	0.000000e+00
19	Road_surface	129.416271	7.899759e-29
20	Surface_condition	683.618300	3.582385e-149
21	Weather	715.145637	5.106202e-156
22	Natural_lighting	505.973206	1.346835e-110
23	Signals_operation	532.837324	1.976265e-116
24	Other_traffic_control	23.932638	6.354679e-06
25	Speed_limit	512.300052	5.694520e-112
26	Road_classification	103.361848	3.591367e-23

Logistic Regression - predict severity of crashes based on features with high chi2 stats: street of crash, type of location, street lighting

```
In [91]: #Preparing data for modelling
X2 = df[['Street_of_crash', 'Street_lighting', 'Type_of_location']]
X2 = pd.get_dummies(X2, columns=['Street_of_crash', 'Street_lighting', 'Type_of_location'])

#Create binary variables to the crash severity variable
df['Crash_severe'] = df['Degree_of_crash'].map({'Fatal':1, 'Injury':1, 'Non-casualty (towaway)':0})
y2 = df[['Crash_severe']]
print(y2)

#Splitting training and testing dataset
from sklearn.model_selection import train_test_split # Import train_test_split function

X2_train, X2_test, y2_train, y2_test = train_test_split(X2, y2, test_size=0.3, random_state=1) # 70% training and 30% test
```

```

      Crash_severe
0                1
1                1
2                1
3                1
4                1
...             ...
101713           1
101714           1
101715           0
101716           0
101717           1

[101718 rows x 1 columns]
```

```
In [92]: pip install -U imbalanced-learn #Import package for undersampling
```

Note: you may need to restart the kernel to use updated packages.

ERROR: Invalid requirement: '#Import'

```
In [93]: from imblearn.under_sampling import RandomUnderSampler # Import the necessary libraries
rus = RandomUnderSampler(random_state=42, sampling_strategy = 'majority')
# Balancing the data
X2_resampled, y2_resampled = rus.fit_resample(X2_train, y2_train)
#Build model
r_accident_lr = LogisticRegression(solver = 'liblinear', random_state = 42, C = 10)
```

```
In [94]: #Step 3: Fit the model with data
r_accident_lr.fit(X2_train, y2_train)

#Step 4: Predicting
y2_pred = r_accident_lr.predict(X2_test)

print(pd.DataFrame(y2_pred).value_counts())

#Step 5 - Evaluate test the model
accuracy = r_accident_lr.score(X2_test, y2_test)
print(accuracy)

#Confusion matrix
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

```

def confusion_matrix_visual(y2_true, y2_pred, labels):
    cm = confusion_matrix(y2_true, y2_pred)

    plt.figure(figsize=(6, 3))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.title('Confusion Matrix')
    plt.show()

confusion_matrix_visual(y2_test, y2_pred, ['low', 'high'])

#Precision report
from sklearn.metrics import classification_report

print("Report : ",
      classification_report(y2_test, y2_pred))

# ROC Curves
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

fpr, tpr, thresholds = roc_curve(y2_test, r_accident_lr.predict_proba(X2_test)[: , 1])
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='red', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()

```


C:\Users\maikh\anaconda3\lib\site-packages\sklearn\utils\validation.py:993: DataConversionWarning:

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```

1    23129
0     7387
dtype: int64
0.7166404509109975

```



```

Report :
          precision    recall  f1-score   support

     0       0.56      0.43      0.49       9462
     1       0.77      0.84      0.80      21054

 accuracy          0.66
 macro avg          0.64
 weighted avg          0.70

```

